



Source

OBJECT DETECTION

Creating a Weapon Detector in 5 simple steps

Object detection using mask-RCNN on custom Dataset



Rahul Agarwal [Follow](#)

Oct 10, 2019 · 7 min read ★

Object Detection is a helpful tool to have in your coding repository.

It forms the backbone of many fantastic industrial applications. Some of them being self-driving cars, medical imaging and face detection.

In my last post on Object detection, I talked about how Object detection models evolved.

But what good is theory, if we can't implement it?

This post is about implementing and getting an object detector on our custom dataset of weapons.

The problem we will specifically solve today is that of Instance Segmentation using Mask-RCNN.

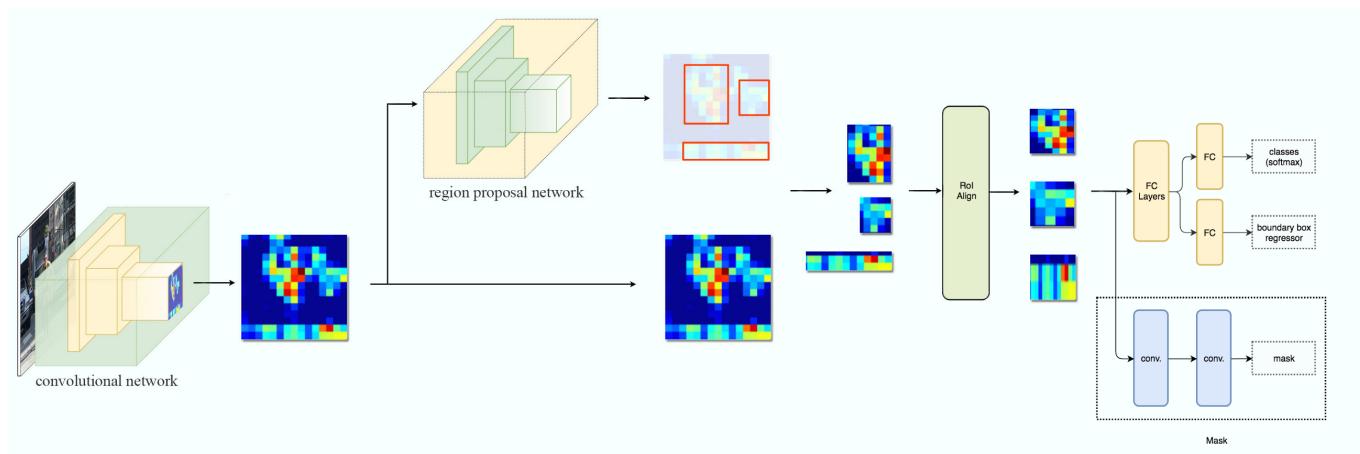
• • •

Instance Segmentation

Can we create masks for each object in the image? Specifically something like:



The most common way to solve this problem is by using Mask-RCNN. The architecture of Mask-RCNN looks like below:



Essentially, it comprises of:

- A backbone network like resnet50/resnet101
- A Region Proposal network
- ROI-Align layers
- Two output layers — one to predict masks and one to predict class and bounding box.

There is a lot more to it. If you want to learn more about the theory, read my last post.

Demystifying Object Detection and Instance Segmentation for Data Scientists

Easy Explanation!!! I tried

towardsdatascience.com

This post is mostly going to be about the code.

• • •

1. Creating your Custom Dataset for Instance Segmentation



Our Dataset

The use case we will be working on is a weapon detector. A weapon detector is something that can be used in conjunction with street cameras as well as CCTV's to fight crime. So it is pretty nifty.

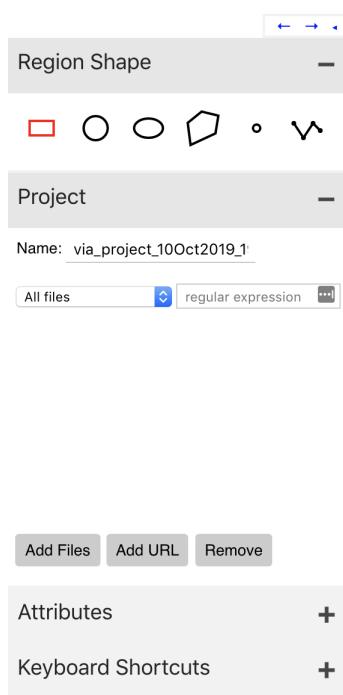
So, I started with downloading 40 images each of guns and swords from the open image dataset and annotated them using the VIA tool. Now setting up the annotation project in VIA is pretty important, so I will try to explain it step by step.

1. Set up VIA

VIA is an annotation tool, using which you can annotate images both bounding boxes as well as masks. I found it as one of the best tools to do annotation as it is online and runs in the browser itself.

To use it, open <http://www.robots.ox.ac.uk/~vgg/software/via/via.html>

You will see a page like:



- To start annotation, [select images](#) (or, add images from [URL](#) or [absolute path](#)) and draw regions
- Use [attribute editor](#) to define attributes (e.g. name) and [annotation editor](#) to describe each region (e.g. cat) using these attributes.
- Remember to [save](#) your project before closing this application so that you can [load](#) it later to continue annotation.
- For help, see the [Getting Started](#) page and pre-loaded demo: [image annotation](#) and [face annotation](#).

The next thing we want to do is to add the different class names in the region_attributes. Here I have added ‘gun’ and ‘sword’ as per our use case as these are the two distinct targets I want to annotate.

The screenshot shows the "Attributes" section of the application. At the top, the title "Attributes" is displayed. Below it, there are two tabs: "Region Attributes" (which is underlined, indicating it is selected) and "File Attributes".

Under the "Region Attributes" tab, there is a search bar with the placeholder "name" and a "..." button. To the right of the search bar are three buttons: a plus sign (+), a minus sign (-), and a dropdown arrow.

Below the search bar, there is a list item with the "name" placeholder. This item has a dropdown arrow to its right.

At the bottom, there is a detailed view for the "name" attribute. It includes fields for "Name" (containing "name") and "Desc." (an empty text area).

I type

checkbox

id	description	def.
gun		<input type="checkbox"/>
sword		<input type="checkbox"/>
Add new option id		
Toggle Annotation Editor		

2. Annotate the Images

I have kept all the files in the folder `data`. Next step is to add the files we want to annotate. We can add files in the `data` folder using the “Add Files” button in the VIA tool. And start annotating along with labels as shown below after selecting the polyline tool.

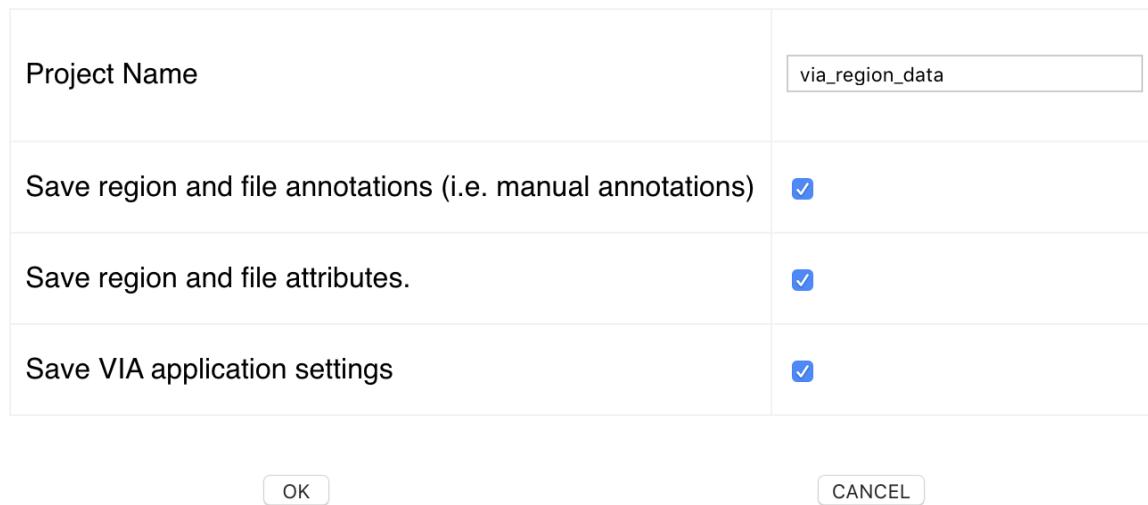


Click, Click, Enter, Escape, Select

3. Download the annotation file

Click on `save project` on the top menu of the VIA tool.

Save Project



Save file as `via_region_data.json` by changing the project name field. This will save the annotations in COCO format.

4. Set up the data directory structure

We will need to set up the data directories first so that we can do object detection. In the code below, I am creating a directory structure that is required for the model that we are going to use.

```

1  from random import random
2  import os
3  from glob import glob
4  import json
5  # Path to your images
6  image_paths = glob("data/*")
7  #Path to your annotations from VIA tool
8  annotation_file = 'via_region_data.json'
9  #clean up the annotations a little
10 annotations = json.load(open(annotation_file))
11 cleaned_annotations = {}
12 for k,v in annotations['via_im_metadata'].items():

```

```
13     cleaned_annotations[v['filename']] = v
14 # create train and validation directories
15 ! mkdir prodata
16 ! mkdir prodata/val
17 ! mkdir prodata/train
18 train_annotations = {}
19 valid_annotations = {}
20 # 20% of images in validation folder
21 for img in image_paths:
22     # Image goes to Validation folder
23     if random()<0.2:
24         os.system("cp "+ img + " prodata/val/")
25         img = img.split("/")[-1]
26         valid_annotations[img] = cleaned_annotations[img]
27     else:
28         os.system("cp "+ img + " prodata/train/")
29         img = img.split("/")[-1]
30         train_annotations[img] = cleaned_annotations[img]
31 # put different annotations in different folders
32 with open('prodata/val/via_region_data.json', 'w') as fp:
33     json.dump(valid_annotations, fp)
34 with open('prodata/train/via_region_data.json', 'w') as fp:
35     json.dump(train_annotations, fp)
```

directory.py hosted with ❤ by GitHub

[view raw](#)

After running the above code, we will get the data in the below folder structure:

- prodata
 - train
 - img1.jpg
 - img2.jpg
 - via_region_data.json
 - val
 - img3.jpg
 - img4.jpg
 - via_region_data.json
- • •

2. Setup the Coding Environment

We will use the code from the `matterport/Mask_RCNN` GitHub repository. You can start by cloning the repository and installing the required libraries.

```
git clone https://github.com/matterport/Mask_RCNN
cd Mask_RCNN
pip install -r requirements.txt
```

Once we are done with installing the dependencies and cloning the repo, we can start with implementing our project.

We make a copy of the `samples/balloon` directory in `Mask_RCNN` folder and create a `samples/guns_and_swords` directory where we will continue our work:

```
cp -r samples/balloon samples/guns_and_swords
```

Setting up the Code



```
...removed for brevity...
```

The image features a composite of two visual elements. On the left, there is a block of heavily obfuscated JavaScript code. On the right, there is a stylized, glowing blue and purple illustration of a circuit board. A human eye is positioned in the center of the circuit board, looking directly forward. The overall theme is the integration of human vision with advanced computer technology and data processing.



Yes. We are doing AI

We start by renaming and changing balloon.py in the ***samples/guns_and_swords directory to gns.py***. The balloon.py file right now trains for one target. I have extended it to use multiple targets. In this file, we change:

1. balloonconfig to gnsConfig
2. BalloonDataset to gnsDataset : We changed some code here to get the target names from our annotation data and also give multiple targets.
3. And some changes in the train function

Showing only the changed gnsConfig here to get you an idea. You can take a look at the whole gns.py code here.

```

1  class gnsConfig(Config):
2      """Configuration for training on the toy dataset.
3      Derives from the base Config class and overrides some values.
4      """
5
6      # Give the configuration a recognizable name
7      NAME = "gns"
8
9      # We use a GPU with 16GB memory, which can fit three images.
10     # Adjust down if you use a smaller GPU.
11     IMAGES_PER_GPU = 3
12
13     # Number of classes (including background)
14     NUM_CLASSES = 1 + 2 # Background + sword + gun
15
16     # Number of training steps per epoch

```

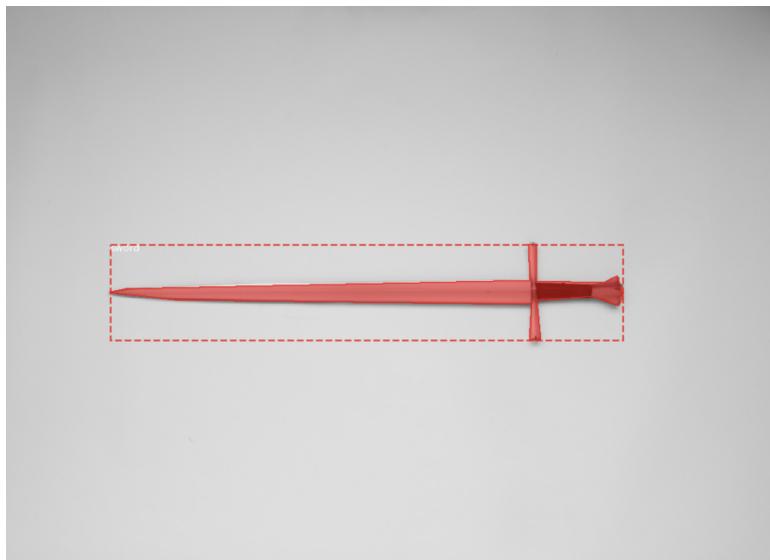
gnsConfig.py hosted with ❤ by GitHub

[view raw](#)

• • •

3. Visualizing Images and Masks

Once we are done with changing the gns.py file, we can visualize our masks and images. You can do simply by following this Visualize Dataset.ipynb notebook.



• • •

4. Train the MaskRCNN Model with Transfer Learning

To train the maskRCNN model, on the Guns and Swords dataset, we need to run one of the following commands on the command line based on if we want to initialise our model with COCO weights or imagenet weights:

```
# Train a new model starting from pre-trained COCO weights
python3 gns.py train --dataset=/path/to/dataset --weights=coco
```

```
# Resume training a model that you had trained earlier  
python3 gns.py train --dataset=/path/to/dataset --weights=last  
  
# Train a new model starting from ImageNet weights  
python3 gns.py train --dataset=/path/to/dataset --weights=imagenet
```

The command with `weights=last` will resume training from the last epoch. The weights are going to be saved in the `logs` directory in the `Mask_RCNN` folder.

This is how the loss looks after our final epoch.

Visualize the losses using Tensorboard

You can take advantage of tensorboard to visualise how your network is performing. Just run:

```
tensorboard --logdir  
~/objectDetection/Mask_RCNN/logs/gns20191010T1234
```

You can get the tensorboard at

<https://localhost:6006>

Here is how our mask loss looks like:

We can see that the validation loss is performing pretty abruptly. This is expected as we only have kept 20 images in the validation set.

• • •

5. Prediction on New Images

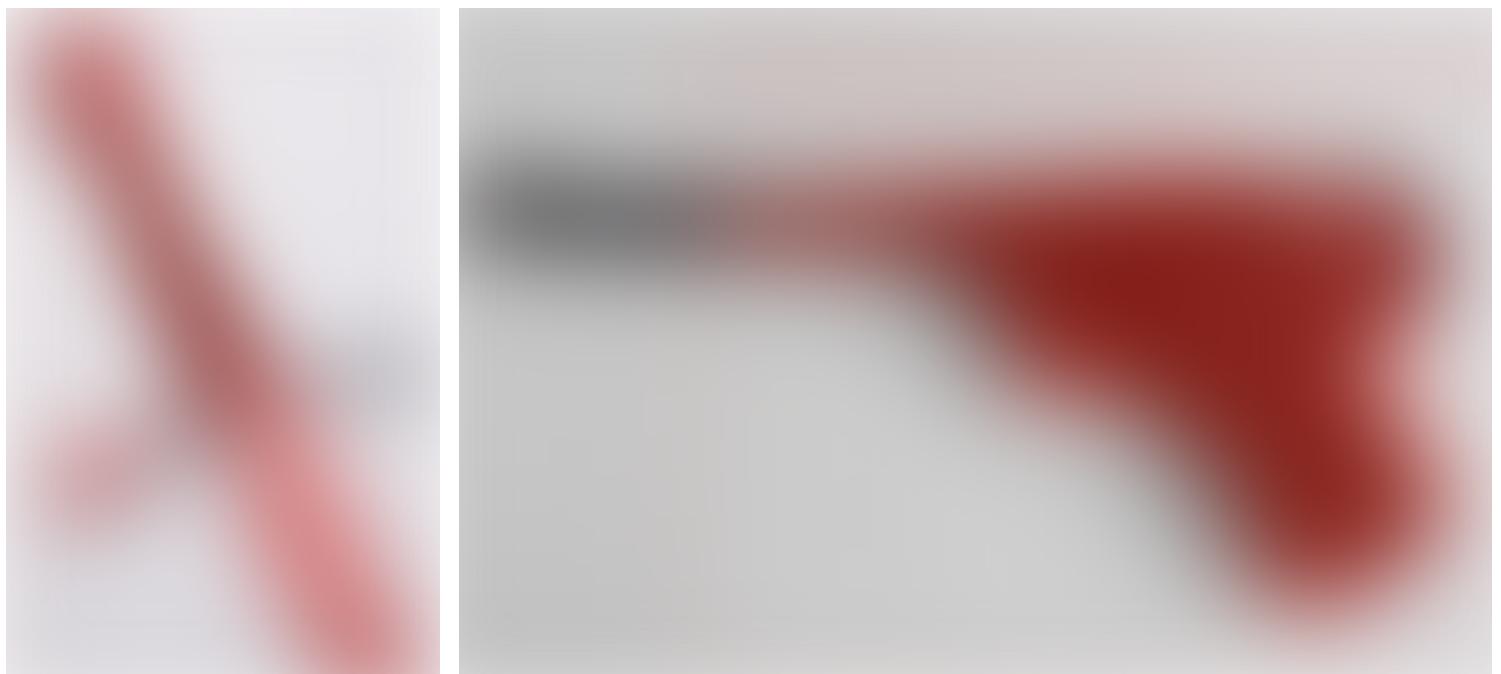
Predicting a new image is also pretty easy. Just follow the `prediction.ipynb` notebook for a minimal example using our trained model. Below is the main part of the code.

```
1 # Function taken from utils.dataset
2 def load_image(image_path):
3     """Load the specified image and return a [H,W,3] Numpy array.
4     """
5     # Load image
6     image = skimage.io.imread(image_path)
7     # If grayscale. Convert to RGB for consistency.
8     if image.ndim != 3:
9         image = skimage.color.gray2rgb(image)
10    # If has an alpha channel, remove it for consistency
11    if image.shape[-1] == 4:
12        image = image[...,:3]
13    return image
14 # path to image to be predicted
15 image = load_image("../data/2c8ce42709516c79.jpg")
16 # Run object detection
17 results = model.detect([image], verbose=1)
18 # Display results
19 ax = get_ax(1)
20 r = results[0]
21 a = visualize.display_instances(image, r['rois'], r['masks'], r['class_ids'], dataset.class_names,
22                                 title="Predictions")
```

[prediction.py](#) hosted with ❤ by GitHub

[view raw](#)

This is how the result looks for some images in the validation set:



• • •

Improvements

The results don't look very promising and leave a lot to be desired, but that is to be expected because of very less training data(60 images). One can try to do the below things to improve the model performance for this weapon detector.

1. We just trained on 60 images due to time constraints. While we used transfer learning the data is still too less — Annotate more data.
2. Train for more epochs and longer time. See how validation loss and training loss looks like.
3. Change hyperparameters in the `mrcnn/config` file in the `Mask_RCNN` directory. For information on what these hyperparameters mean, take a look at my previous post. The main ones you can look at:

```
# if you want to provide different weights to different losses
LOSS_WEIGHTS ={'rpn_class_loss': 1.0, 'rpn_bbox_loss': 1.0,
```

```
'mrcnn_class_loss': 1.0, 'mrcnn_bbox_loss': 1.0, 'mrcnn_mask_loss':  
1.0}  
  
# Length of square anchor side in pixels  
RPN_ANCHOR_SCALES = (32, 64, 128, 256, 512)  
  
# Ratios of anchors at each cell (width/height)  
# A value of 1 represents a square anchor, and 0.5 is a wide anchor  
RPN_ANCHOR_RATIOS = [0.5, 1, 2]
```

• • •

Conclusion



Photo by Christopher Gower on Unsplash

In this post, I talked about how to implement Instance segmentation using Mask-RCNN for a custom dataset.

I tried to make the coding part as simple as possible and hope you find the code useful. In the next part of this post, I will deploy this model using a web app. So stay tuned.

You can download the annotated weapons data as well as the code at [Github](#).

• • •

If you want to know more about various ***Object Detection techniques, motion estimation, object tracking in video etc.***, I would like to recommend this awesome course on Deep Learning in Computer Vision in the Advanced machine learning specialization.

• • •

Thanks for the read. I am going to be writing more beginner-friendly posts in the future too. Follow me up at [Medium](#) or [Subscribe to my blog](#) to be informed about them. As always, I welcome feedback and constructive criticism and can be reached on Twitter @mlwhiz.

Also, a small disclaimer — There might be some affiliate links in this post to relevant resources as sharing knowledge is never a bad idea.

[Machine Learning](#) [Artificial Intelligence](#) [Programming](#) [Data Science](#) [Towards Data Science](#)

[About](#) [Help](#) [Legal](#)